

005 -- Backtest and Validation

Author: Wayne Kirk Schmidt
Email: wayne.kirk.schmidt@gmail.com

Purpose

`005_backtest.ipynb` is the fifth and final analytical stage of the cryptocurrency statistical arbitrage research pipeline.

This notebook evaluates the leader-follower signal identified in Stage 4 by simulating trades under realistic execution assumptions, measuring performance, and testing statistical robustness.

The goal of this stage is **validation**: to determine whether the edge observed in Stage 4 survives realistic transaction costs, execution delays, and out-of-sample time periods.

Inputs

From `output/004_strategy/` and `output/002_enrich/`:

File	Description
<code>event_response_expanded.pkl</code>	Event-level leader-follower response dataset with per-lag returns
<code>returns_full.pkl</code>	Daily return series for all assets
<code>price_wide.pkl</code>	Wide-format price series (for diagnostics)

Outputs

Written to `output/005_backtest/`:

File	Description
<code>trade_df.pkl</code>	Full trade-level dataset
<code>daily_returns.pkl</code>	Aggregated daily return series
<code>equity_curve.pkl</code>	Cumulative equity curve
<code>performance_summary.pkl</code>	Summary performance metrics
<code>grid_results.pkl</code>	Stress-grid results across all delay/cost scenarios
<code>walkforward_results.pkl</code>	Walk-forward fold performance

Tasks

1. **Signal construction** -- rebuild the filtered and strong event sets from the expanded event response data
2. **Trade construction** -- map signals to directional trades using t+n forward returns
3. **Trade simulation** -- run across a grid of execution delays (0-3 days) and transaction costs (20-40 bps)
4. **Performance evaluation** -- Sharpe ratio, drawdown, hit rate
5. **Statistical significance** -- t-test on daily returns; Sharpe CI via Lo (2002)
6. **Walk-forward robustness** -- performance across 3 non-overlapping time folds

Pipeline Position

```
001_download -> data acquisition
002_enrich   -> feature engineering
003_analysis -> shock detection
004_strategy -> signal construction
005_backtest <- you are here
```

Design Principles

- **Deterministic** -- operates exclusively on persisted artifacts; no recomputation of prior stages
- **No forward-looking assumptions** -- all results derived from information available at time of event
- **Full artifact persistence** -- every output is saved and registered in the manifest

1. Imports and Environment Setup

Standard library imports for data manipulation, statistics, file I/O, and visualization.

```
In [1]: import pandas as pd
import numpy as np
import pickle
from datetime import datetime, UTC
import math
from pathlib import Path
import matplotlib.pyplot as plt
from scipy import stats
```

2. Environment Configuration

Defines the asset universe, pipeline parameters, and output directories for this stage.

```
In [2]: startdate = "2023-01-01"
trading_days = 252
frequency = "1d"

universe = [
    "BTCUSD", # Bitcoin
    "ETHUSD", # Ethereum
    "BNBUSDT", # Binance Coin
    "SOLUSD", # Solana
    "XRPUSD", # Ripple
    "ADAUSD", # Cardano
    "DOGEUSD", # Dogecoin
    "AVAXUSD", # Avalanche
    "LTCUSD" # Litecoin
]

execution_delay = [0, 1, 2, 3]
execution_cost_bps = [20, 30, 40]

stage_label = "005_backtest"

OUTPUT_ROOT = Path("../output")
OUTPUT_ROOT.mkdir(parents=True, exist_ok=True)

MANIFEST_FILE = OUTPUT_ROOT / "manifest.pkl"

DOWNLOAD_DIR = OUTPUT_ROOT / "001_download"
DOWNLOAD_DIR.mkdir(parents=True, exist_ok=True)

ENRICH_DIR = OUTPUT_ROOT / "002_enrich"
ENRICH_DIR.mkdir(parents=True, exist_ok=True)

ANALYSIS_DIR = OUTPUT_ROOT / "003_analysis"
ANALYSIS_DIR.mkdir(parents=True, exist_ok=True)

STRATEGY_DIR = OUTPUT_ROOT / "004_strategy"
STRATEGY_DIR.mkdir(parents=True, exist_ok=True)

BACKTEST_DIR = OUTPUT_ROOT / "005_backtest"
BACKTEST_DIR.mkdir(parents=True, exist_ok=True)

inspection_window = 20

sigma_threshold = 3

observation_window_length = 10
observation_window = range(1, observation_window_length + 1)

holding_period = 1
```

3. Load or Initialize the Manifest

Loads the shared pipeline manifest and initializes this stage's entry if not already present.

```
In [3]: if MANIFEST_FILE.exists():
        manifest = pd.read_pickle(MANIFEST_FILE)
    else:
        manifest = {}

    manifest.setdefault(stage_label, {})
```

```
Out[3]: {}
```

4. Load Required Artifacts from Prior Stages

Loads the event response dataset from Stage 4 and the enriched return/price data from Stage 2. Raises a descriptive error if any required file is missing.

```
In [4]: required_files = {
    "event_response_expanded": (
        STRATEGY_DIR / "event_response_expanded.pkl"
    ),
    "returns_full": (
        ENRICH_DIR / "returns_full.pkl"
    ),
    "price_wide": (
        ENRICH_DIR / "price_wide.pkl"
    ),
}

data = {}

for name, path in required_files.items():
    if not path.exists():
        raise FileNotFoundError(f"Missing required file: {path}")
    data[name] = pd.read_pickle(path)

event_response_expanded = data["event_response_expanded"]
returns_full = data["returns_full"]
price_wide = data["price_wide"]

print("Loaded artifacts:\n")

print("event_response_expanded:", event_response_expanded.shape)
print(event_response_expanded.columns.tolist(), "\n")

print("returns_full:", returns_full.shape)
print("price_wide:", price_wide.shape)
```

loaded artifacts:

```
event_response_expanded: (2320, 14)
['event_date', 'reference_coin', 'target_coin', 'sigma', 'lag_1', 'lag_2', 'lag_3', 'lag_4', 'lag_5', 'lag_6', 'lag_7', 'lag_8', 'lag_9', 'lag_10']

returns_full: (1047, 9)
price_wide: (1242, 9)
```

5. Standardize Column Names

Renames columns to the convention used throughout this notebook (`r_coin` , `f_coin` , `date`) for consistency with the trade construction logic below.

```
In [5]: event_full_df = event_response_expanded.copy()
event_full_df = event_full_df.rename(columns={
    "reference_coin": "r_coin",
    "target_coin": "f_coin",
    "event_date": "date"
})

print("event_full_df:", event_full_df.shape)
print(event_full_df.columns.tolist())

print("\nSample:")
print(event_full_df.head())
```

```
event_full_df: (2320, 14)
['date', 'r_coin', 'f_coin', 'sigma', 'lag_1', 'lag_2', 'lag_3', 'lag_4', 'lag_5', 'lag_6', 'lag_7', 'lag_8', 'lag_9', 'lag_10']
```

Sample:

	date	r_coin	f_coin	sigma	lag_1	lag_2	\
0	2023-08-03 00:00:00+00:00	LTCUSD	ADAUSD	-3	0.004098	-0.001020	
1	2023-08-03 00:00:00+00:00	LTCUSD	AVAXUSD	-3	-0.003213	0.002417	
2	2023-08-03 00:00:00+00:00	LTCUSD	BNBUSD	-3	0.001663	0.007884	
3	2023-08-03 00:00:00+00:00	LTCUSD	BTCUSD	-3	-0.002569	-0.001855	
4	2023-08-03 00:00:00+00:00	LTCUSD	DOGEUSD	-3	-0.001764	0.028688	

	lag_3	lag_4	lag_5	lag_6	lag_7	lag_8	lag_9	\
0	-0.006129	-0.003768	0.023384	0.011089	-0.014955	-0.009447	-0.006131	
1	0.011254	-0.011924	0.021722	-0.004724	-0.011867	-0.004804	-0.004023	
2	0.000412	-0.004938	0.014475	-0.006115	-0.010254	-0.005387	0.000417	
3	0.000888	0.003721	0.019690	-0.006716	-0.004086	-0.000918	-0.000073	
4	-0.017182	-0.011565	0.020952	0.004797	0.005570	-0.001846	0.013742	

	lag_10
0	-0.006511
1	-0.008078
2	0.000416
3	-0.003871
4	-0.024374

6. Build and Filter Event Sets by Signal Direction

Constructs three event sets based on directional consistency between the reference coin's shock and the follower asset's forward returns:

- **Full** -- all events where $\sigma \geq 3$ or ≤ -3
- **Filtered** -- events where `lag_1` direction matches the reference shock direction
- **Strong** -- events where `lag_1`, `lag_2`, and `lag_3` all align with the reference shock direction

The filtered set is used as the primary signal throughout the backtest.

```
In [6]: df = event_response_expanded.copy()

df["r_sign"] = np.sign(df["sigma"])

df["f_sign0"] = np.sign(df["lag_1"])
df["f_sign1"] = np.sign(df["lag_2"])
df["f_sign2"] = np.sign(df["lag_3"])

df = df.rename(columns={
    "reference_coin": "r_coin",
    "target_coin": "f_coin",
    "event_date": "date"
})

event_full_df = df

event_filtered_df = event_full_df[
    event_full_df["r_sign"] == event_full_df["f_sign0"]
]

event_strong_df = event_full_df[
    (event_full_df["r_sign"] == event_full_df["f_sign0"]) &
    (event_full_df["r_sign"] == event_full_df["f_sign1"]) &
    (event_full_df["r_sign"] == event_full_df["f_sign2"])
]

print("columns:")
print(event_full_df.columns.tolist())

print("\ncounts:")
print("full:", len(event_full_df))
print("filtered:", len(event_filtered_df))
print("strong:", len(event_strong_df))

print("\nsample:")
print(event_full_df[[
    "date", "r_coin", "f_coin",
    "sigma", "r_sign",
    "f_sign0", "f_sign1", "f_sign2"
]].head())
```

```

columns:
['date', 'r_coin', 'f_coin', 'sigma', 'lag_1', 'lag_2', 'lag_3', 'lag_4', 'lag_5', 'lag_6', 'lag_7', 'lag_8', 'lag_9', 'lag_10', 'r_sign', 'f_sign0', 'f_sign1', 'f_sign2']

counts:
full: 2320
filtered: 1077
strong: 189

sample:
      date    r_coin    f_coin  sigma  r_sign  f_sign0 \
0 2023-08-03 00:00:00+00:00 LTCUSDT  ADAUSDT    -3    -1    1.0
1 2023-08-03 00:00:00+00:00 LTCUSDT  AVAXUSDT    -3    -1   -1.0
2 2023-08-03 00:00:00+00:00 LTCUSDT  BNBUSDT    -3    -1    1.0
3 2023-08-03 00:00:00+00:00 LTCUSDT  BTCUSDT    -3    -1   -1.0
4 2023-08-03 00:00:00+00:00 LTCUSDT  DOGEUSDT    -3    -1   -1.0

      f_sign1  f_sign2
0      -1.0      -1.0
1       1.0       1.0
2       1.0       1.0
3      -1.0       1.0
4       1.0      -1.0

```

7. Select the Active Event Set

Sets `ACTIVE_SET` to select which event population is used for all downstream trade construction and simulation. Default is `"filtered"` -- the directionally consistent subset.

```
In [7]: ACTIVE_SET = "filtered" # options: "full", "filtered", "strong"
```

```

if ACTIVE_SET == "full":
    events = event_full_df.copy()
elif ACTIVE_SET == "filtered":
    events = event_filtered_df.copy()
elif ACTIVE_SET == "strong":
    events = event_strong_df.copy()
else:
    raise ValueError("Invalid ACTIVE_SET")

print("active set:", ACTIVE_SET)
print("rows:", len(events))
print("columns:", events.columns.tolist())
print("\nsample:")
print(events.head())

```

```

active set: filtered
rows: 1077
columns: ['date', 'r_coin', 'f_coin', 'sigma', 'lag_1', 'lag_2', 'lag_3', 'lag_4', 'lag_5', 'lag_6', 'lag_7', 'lag_8', 'lag_9', 'lag_10', 'r_sign', 'f_sign0', 'f_sign1', 'f_sign2']

sample:
      date    r_coin    f_coin  sigma    lag_1    lag_2 \
1 2023-08-03 00:00:00+00:00 LTCUSDT  AVAXUSDT    -3 -0.003213  0.002417
3 2023-08-03 00:00:00+00:00 LTCUSDT  BTCUSDT    -3 -0.002569 -0.001855
4 2023-08-03 00:00:00+00:00 LTCUSDT  DOGEUSDT    -3 -0.001764  0.028688
5 2023-08-03 00:00:00+00:00 LTCUSDT  ETHUSDT    -3 -0.002887  0.003103
7 2023-08-03 00:00:00+00:00 LTCUSDT  XRPUSDT    -3 -0.045755 -0.010847

      lag_3    lag_4    lag_5    lag_6    lag_7    lag_8    lag_9 \
1  0.011254 -0.011924  0.021722 -0.004724 -0.011867 -0.004804 -0.004023
3  0.000888  0.003721  0.019690 -0.006716 -0.004086 -0.000918 -0.000073
4 -0.017182 -0.011565  0.020952  0.004797  0.005570 -0.001846  0.013742
5 -0.003518 -0.000372  0.015017 -0.000275 -0.001519 -0.002471  0.000654
7 -0.010966  0.001125  0.029053  0.003900 -0.014605 -0.001734 -0.008687

      lag_10  r_sign  f_sign0  f_sign1  f_sign2
1 -0.008078    -1    -1.0     1.0     1.0
3 -0.003871    -1    -1.0    -1.0     1.0
4 -0.024374    -1    -1.0     1.0    -1.0
5 -0.005146    -1    -1.0     1.0    -1.0
7 -0.001434    -1    -1.0    -1.0    -1.0

```

8. Trade Mapping

Converts the active event set into a structured trade representation. Each trade captures:

- Entry return: `lag_1` (t+1 forward return)
- Exit return: `lag_2` (t+2 forward return)
- Direction: sign of the reference coin's shock
- PnL: direction × (entry return + exit return)

```
In [8]: trades = events.copy()

trades["entry_date"] = trades["date"]
trades["exit_date"] = trades["date"]

trades["ret_entry"] = trades["lag_1"] # t+1
trades["ret_exit"] = trades["lag_2"] # t+2

trades["direction"] = trades["r_sign"]

trades["pnl"] = trades["direction"] * (trades["ret_entry"] + trades["ret_exit"])

trade_df = trades[[
    "date",
    "r_coin",
    "f_coin",
    "direction",
    "ret_entry",
    "ret_exit",
    "pnl"
]]

```

```

]]).copy()

print("trade_df shape:", trade_df.shape)
print("\nsample:")
print(trade_df.head())

```

trade_df shape: (1077, 7)

sample:

```

      date    r_coin    f_coin direction net_entry \
1 2023-08-03 00:00:00+00:00 LTCUSDT  AVAXUSDT      -1 -0.003213
3 2023-08-03 00:00:00+00:00 LTCUSDT   BTCUSDT      -1 -0.002569
4 2023-08-03 00:00:00+00:00 LTCUSDT  DOGEUSDT      -1 -0.001764
5 2023-08-03 00:00:00+00:00 LTCUSDT   ETHUSD      -1 -0.002887
7 2023-08-03 00:00:00+00:00 LTCUSDT   XRPUSD      -1 -0.045755

      net_exit      pnl
1  0.002417  0.000795
3 -0.001855  0.004424
4  0.028688 -0.026924
5  0.003103 -0.000216
7 -0.010847  0.056602

```

9. Daily Aggregation and Gross Performance

Aggregates trade PnL by date (one trade per day, unbiased baseline) to produce a daily return series and cumulative equity curve. Computes gross performance metrics before transaction costs are applied.

```

In [9]: # Daily aggregation -- equal-weighted mean of all trade PnLs per day
# NOTE: this is a GROSS, PRE-COST, UNADJUSTED metric.
# It reflects raw signal strength only.
# For net performance and proper Sharpe, see the stress grid below.
daily_returns = (
    trade_df
    .groupby("date")["pnl"]
    .mean()
    .sort_index()
)

equity_curve = (1 + daily_returns).cumprod()

print("daily_returns shape:", daily_returns.shape)
print("NOTE: gross pre-cost unadjusted -- see stress grid for net results")
print()
print("Daily return distribution:")
print(daily_returns.describe().round(4))

```

daily_returns shape: (86,)

NOTE: gross pre-cost unadjusted -- see stress grid for net results

Daily return distribution:

```

count    86.0000
mean      0.0324
std       0.0486
min      -0.0850
25%       0.0028
50%       0.0245
75%       0.0647
max       0.2110

```

Name: pnl, dtype: float64

10. Net Performance Metrics

Computes the final performance summary over the full trade set including total return, annualized volatility, and Sharpe ratio. This serves as the baseline for the stress-grid evaluation below.

```

In [10]: # GROSS UNADJUSTED PERFORMANCE SUMMARY
# This cell computes trade-level metrics WITHOUT:
#   - transaction costs
#   - execution delay
#   - daily aggregation
# These numbers are intentionally inflated and are shown only to
# illustrate raw signal strength. Do NOT cite these as strategy performance.
# See the stress grid (Section 11) for all net performance metrics.

trade_returns = trade_df["pnl"].dropna()

total_return_gross = trade_returns.sum()
sharpe_gross = (
    trade_returns.mean() / trade_returns.std()
) * np.sqrt(trading_days)

print("=== GROSS SIGNAL STRENGTH (pre-cost, trade-level) ===")
print(f" num_trades:    {len(trade_returns)}")
print(f" mean per trade:{trade_returns.mean():>+.4f}")
print(f" gross Sharpe:   {sharpe_gross:.2f} (inflated -- not meaningful)")
print()
print("--> See Section 11 (stress grid) for actual net performance.")

performance_summary = pd.Series({
    "total_return": total_return_gross,
    "volatility":   trade_returns.std() * np.sqrt(trading_days),
    "sharpe":       sharpe_gross,
    "num_trades":   float(len(trade_returns)),
})

```

=== GROSS SIGNAL STRENGTH (pre-cost, trade-level) ===

```

num_trades:    1077
mean per trade:+0.0326
gross Sharpe:   8.95 (inflated -- not meaningful)

```

--> See Section 11 (stress grid) for actual net performance.

```
In [11]: # Helper: reconstruct daily net PnL for any delay/cost scenario
# Used by sections 11, 16, and 17.
def scenario_daily(df, delay, cost_bps):
    """Return daily net PnL series for given execution delay and cost."""
    cost = cost_bps / 10000.0
    sc = df.copy()
    entry_col = f"lag_{1 + delay}"
    exit_col = f"lag_{2 + delay}"
    if entry_col not in sc.columns or exit_col not in sc.columns:
        return pd.Series(dtype=float)
    sc["direction"] = sc["r_sign"]
    leg1 = sc[["date", "f_coin", "direction", entry_col]].copy()
    leg1["date"] += pd.Timedelta(days=1)
    leg1["net_pnl"] = leg1["direction"] * leg1[entry_col] - cost / 2
    leg2 = sc[["date", "f_coin", "direction", exit_col]].copy()
    leg2["date"] += pd.Timedelta(days=2)
    leg2["net_pnl"] = leg2["direction"] * leg2[exit_col] - cost / 2
    expanded = pd.concat([leg1, leg2], ignore_index=True)
    expanded = (
        expanded
        .sort_values(["date", "f_coin"])
        .groupby("date")
        .head(1)
    )
    return expanded.set_index("date")["net_pnl"].sort_index()
```

11. Run the Execution Stress Grid

Simulates the strategy across all combinations of:

- **Execution delay:** 0, 1, 2, 3 days after the shock event
- **Transaction cost:** 20, 30, 40 basis points (round-trip)

For each scenario, computes total return, annualized volatility, Sharpe ratio, and maximum drawdown. Results are collected for comparison in the next section.

```
In [12]: results = []

for delay in execution_delay:
    for cost_bps in execution_cost_bps:

        cost = cost_bps / 10000.0

        scenario_df = events.copy()

        entry_col = f"lag_{1 + delay}"
        exit_col = f"lag_{2 + delay}"

        if entry_col not in scenario_df.columns or exit_col not in scenario_df.columns:
            continue

        scenario_df["ret_entry"] = scenario_df[entry_col]
        scenario_df["ret_exit"] = scenario_df[exit_col]

        scenario_df["direction"] = scenario_df["r_sign"]

        # --- PnL construction ---
        leg1 = scenario_df[["date", "f_coin", "direction", "ret_entry"]].copy()
        leg1["date"] = leg1["date"] + pd.Timedelta(days=1)
        leg1["net_pnl"] = leg1["direction"] * leg1["ret_entry"] - cost / 2

        leg2 = scenario_df[["date", "f_coin", "direction", "ret_exit"]].copy()
        leg2["date"] = leg2["date"] + pd.Timedelta(days=2)
        leg2["net_pnl"] = leg2["direction"] * leg2["ret_exit"] - cost / 2

        expanded = pd.concat([leg1, leg2], ignore_index=True)

        # --- keep 1 trade per day (unbiased baseline) ---
        expanded = (
            expanded
            .sort_values(["date", "f_coin"])
            .groupby("date")
            .head(1)
        )

        daily_net = expanded.set_index("date")["net_pnl"].sort_index()

        if len(daily_net) < 2:
            continue

        equity = (1 + daily_net).cumprod()

        total_return = equity.iloc[-1] - 1

        volatility = daily_net.std() * np.sqrt(trading_days)

        sharpe = (
            daily_net.mean() / daily_net.std()
        ) * np.sqrt(trading_days)

        drawdown = equity / equity.cummax() - 1
        max_drawdown = drawdown.min()

        results.append({
            "delay": delay,
            "cost_bps": cost_bps,
            "num_days": len(daily_net),
            "total_return": total_return,
            "volatility": volatility,
            "sharpe": sharpe,
            "max_drawdown": max_drawdown,
```

```
)
print("scenarios calculated:", len(results))

results
```

scenarios calculated: 12

```
Out[12]: [{'delay': 0,
          'cost_bps': 20,
          'num_days': 159,
          'total_return': np.float64(3.027864131810766),
          'volatility': np.float64(0.7501198261238219),
          'sharpe': np.float64(3.3250155789485416),
          'max_drawdown': np.float64(-0.2208451015202808)},
         {'delay': 0,
          'cost_bps': 30,
          'num_days': 159,
          'total_return': np.float64(2.722235617368532),
          'volatility': np.float64(0.7501198261238219),
          'sharpe': np.float64(3.157042415712616),
          'max_drawdown': np.float64(-0.2290875568660451)},
         {'delay': 0,
          'cost_bps': 40,
          'num_days': 159,
          'total_return': np.float64(2.439662678459568),
          'volatility': np.float64(0.7501198261238219),
          'sharpe': np.float64(2.989069252476692),
          'max_drawdown': np.float64(-0.2372469333657511)},
         {'delay': 1,
          'cost_bps': 20,
          'num_days': 159,
          'total_return': np.float64(-0.302712803039035),
          'volatility': np.float64(0.7736524520698632),
          'sharpe': np.float64(-0.3591083171203813),
          'max_drawdown': np.float64(-0.4794192708058075)},
         {'delay': 1,
          'cost_bps': 30,
          'num_days': 159,
          'total_return': np.float64(-0.3561890794945566),
          'volatility': np.float64(0.7736524520698632),
          'sharpe': np.float64(-0.5219721452681423),
          'max_drawdown': np.float64(-0.5092033694139568)},
         {'delay': 1,
          'cost_bps': 40,
          'num_days': 159,
          'total_return': np.float64(-0.40558802776205716),
          'volatility': np.float64(0.7736524520698632),
          'sharpe': np.float64(-0.6848359734159033),
          'max_drawdown': np.float64(-0.5372971917732847)},
         {'delay': 2,
          'cost_bps': 20,
          'num_days': 159,
          'total_return': np.float64(-0.0050216060745049695),
          'volatility': np.float64(0.7510398842678386),
          'sharpe': np.float64(0.3533862928354955),
          'max_drawdown': np.float64(-0.40547181277743216)},
         {'delay': 2,
          'cost_bps': 30,
          'num_days': 159,
          'total_return': np.float64(-0.08115824777084046),
          'volatility': np.float64(0.7510398842678386),
          'sharpe': np.float64(0.18561890439269277),
          'max_drawdown': np.float64(-0.42349992836713246)},
         {'delay': 2,
          'cost_bps': 40,
          'num_days': 159,
          'total_return': np.float64(-0.15150275251772594),
          'volatility': np.float64(0.7510398842678386),
          'sharpe': np.float64(0.017851515949889988),
          'max_drawdown': np.float64(-0.4409900776608766)},
         {'delay': 3,
          'cost_bps': 20,
          'num_days': 159,
          'total_return': np.float64(0.49200977189075834),
          'volatility': np.float64(0.7502328487829448),
          'sharpe': np.float64(1.211209376556388),
          'max_drawdown': np.float64(-0.2796572949078937)},
         {'delay': 3,
          'cost_bps': 30,
          'num_days': 159,
          'total_return': np.float64(0.37811858975026746),
          'volatility': np.float64(0.7502328487829448),
          'sharpe': np.float64(1.0432615184955183),
          'max_drawdown': np.float64(-0.29043414938940837)},
         {'delay': 3,
          'cost_bps': 40,
          'num_days': 159,
          'total_return': np.float64(0.2728705766161641),
          'volatility': np.float64(0.7502328487829448),
          'sharpe': np.float64(0.8753136604346482),
          'max_drawdown': np.float64(-0.3041667613928871)}]
```

12. Stress Grid Summary Table

Displays the full stress grid as a formatted DataFrame sorted by execution delay and cost. This is the primary deliverable of the backtest -- a complete picture of how performance varies across realistic execution assumptions.

```
In [13]: pd.set_option("display.expand_frame_repr", False)
pd.set_option("display.max_columns", None)

stress_grid = (
    pd.DataFrame(results)
```

```

        .sort_values(["delay", "cost_bps"])
        .reset_index(drop=True)
    )

    print("Stress Grid:")
    print(stress_grid)

    print("\nShape:", stress_grid.shape)

```

```

Stress Grid:
   delay  cost_bps  num_days  total_return  volatility  sharpe  max_drawdown
0      0         20       159      3.027864      0.750120  3.325016      -0.220845
1      0         30       159      2.722236      0.750120  3.157042      -0.229088
2      0         40       159      2.439663      0.750120  2.989069      -0.237247
3      1         20       159     -0.302713      0.773652 -0.359108      -0.479419
4      1         30       159     -0.356189      0.773652 -0.521972      -0.509203
5      1         40       159     -0.405588      0.773652 -0.684836      -0.537297
6      2         20       159     -0.005022      0.751040  0.353386      -0.405472
7      2         30       159     -0.081158      0.751040  0.185619      -0.423500
8      2         40       159     -0.151503      0.751040  0.017852      -0.440990
9      3         20       159      0.492010      0.750233  1.211209      -0.279657
10     3         30       159      0.378119      0.750233  1.043262      -0.290434
11     3         40       159      0.272871      0.750233  0.875314      -0.304167

```

Shape: (12, 7)

13. Sharpe Ratio Summary by Scenario

Prints the Sharpe ratio for each (delay, cost) combination in a compact table for quick review.

```
In [14]: print(stress_grid[["delay", "cost_bps", "sharpe"]])
```

```

   delay  cost_bps  sharpe
0      0         20  3.325016
1      0         30  3.157042
2      0         40  2.989069
3      1         20 -0.359108
4      1         30 -0.521972
5      1         40 -0.684836
6      2         20  0.353386
7      2         30  0.185619
8      2         40  0.017852
9      3         20  1.211209
10     3         30  1.043262
11     3         40  0.875314

```

14. Visualize Sharpe Ratio Across the Stress Grid

Plots Sharpe ratio as a function of execution delay for each cost level. A strategy with consistently positive Sharpe across the grid -- and graceful degradation as delay and cost increase -- shows robustness to execution friction.

```

In [15]: pivot_sharpe = stress_grid.pivot(
        index="delay",
        columns="cost_bps",
        values="sharpe"
    )

    plt.figure()

    for col in pivot_sharpe.columns:
        plt.plot(pivot_sharpe.index, pivot_sharpe[col], marker='o', label=f"cost {col} bps")

    plt.xlabel("Execution Delay (days)")
    plt.ylabel("Sharpe Ratio")
    plt.title("Sharpe vs Execution Delay (by Cost)")


    plt.legend()
    plt.grid()

    plt.show()

    print("\n")

    print(daily_net.describe())

```




```
count      159.000000
mean       0.002606
std        0.047260
min        -0.111803
25%        -0.019485
50%         0.002686
75%         0.027145
max         0.216053
Name: net_pnl, dtype: float64
```

15. Persist Backtest Artifacts

Saves all primary backtest outputs to disk and registers each in the manifest with file existence and non-empty size assertions.

```
In [16]: fp = BACKTEST_DIR / "trade_df.pkl"
trade_df.to_pickle(fp)
assert fp.exists(), "trade_df.pkl not written"
assert fp.stat().st_size > 0, "trade_df.pkl is empty"
manifest[stage_label]["trade_df"] = str(fp.relative_to(OUTPUT_ROOT))
print("saved:", fp)

fp = BACKTEST_DIR / "daily_returns.pkl"
daily_returns.to_pickle(fp)
assert fp.exists(), "daily_returns.pkl not written"
assert fp.stat().st_size > 0, "daily_returns.pkl is empty"
manifest[stage_label]["daily_returns"] = str(fp.relative_to(OUTPUT_ROOT))
print("saved:", fp)

fp = BACKTEST_DIR / "equity_curve.pkl"
equity_curve.to_pickle(fp)
assert fp.exists(), "equity_curve.pkl not written"
assert fp.stat().st_size > 0, "equity_curve.pkl is empty"
manifest[stage_label]["equity_curve"] = str(fp.relative_to(OUTPUT_ROOT))
print("saved:", fp)

fp = BACKTEST_DIR / "performance_summary.pkl"
performance_summary.to_pickle(fp)
assert fp.exists(), "performance_summary.pkl not written"
assert fp.stat().st_size > 0, "performance_summary.pkl is empty"
manifest[stage_label]["performance_summary"] = str(fp.relative_to(OUTPUT_ROOT))
print("saved:", fp)

fp = BACKTEST_DIR / "grid_results.pkl"
results_df = pd.DataFrame(results)
results_df.to_pickle(fp)
assert fp.exists(), "grid_results.pkl not written"
assert fp.stat().st_size > 0, "grid_results.pkl is empty"
manifest[stage_label]["grid_results"] = str(fp.relative_to(OUTPUT_ROOT))
print("saved:", fp)
```

```
saved: ..\output\005_backtest\trade_df.pkl
saved: ..\output\005_backtest\daily_returns.pkl
saved: ..\output\005_backtest\equity_curve.pkl
saved: ..\output\005_backtest\performance_summary.pkl
saved: ..\output\005_backtest\grid_results.pkl
```

16. Statistical Significance Tests

We test whether the observed mean return is statistically distinguishable from zero, and compute a confidence interval around the annualized Sharpe ratio.

16a -- One-sample t-test on daily returns (t+0, 20 bps baseline)

Null hypothesis: mean daily return = 0. A p-value below 0.05 provides evidence that the strategy's returns are not explained by chance.

16b -- Sharpe ratio confidence interval (Lo, 2002)

The Lo (2002) formula adjusts the standard error of the Sharpe ratio for non-normality (skewness and excess kurtosis) in the return series -- an important correction for cryptocurrency returns, which exhibit heavy tails. A 95% CI that excludes zero provides evidence the Sharpe is not a sampling artifact.

```
In [17]: from scipy import stats

# --- 16a: t-test on daily returns ---
daily_base = scenario_daily(event_filtered_df, delay=0, cost_bps=20)

t_stat, p_value = stats.ttest_1samp(daily_base.dropna(), popmean=0)
n = len(daily_base)
mean_ret = daily_base.mean()
std_ret = daily_base.std()
se = std_ret / np.sqrt(n)
ci_lo = mean_ret - 1.96 * se
ci_hi = mean_ret + 1.96 * se

print("=== 16a: One-sample t-test -- daily return vs 0 ===")
print(f" n : {n}")
print(f" mean return : {mean_ret:.4f} ({mean_ret*100:.2f}% per day)")
print(f" t-statistic : {t_stat:.3f}")
print(f" p-value : {p_value:.4f}")
print(f" 95% CI : [{ci_lo:.4f}, {ci_hi:.4f}])")
print()
if p_value < 0.05:
    print(" Interpretation: reject H0 at 5% -- mean return is significantly positive.")
else:
    print(" Interpretation: cannot reject H0 at 5% -- insufficient evidence of positive mean.")

print()

# --- 16b: Sharpe CI via Lo (2002) ---
sr = (mean_ret / std_ret) * np.sqrt(trading_days)
skewness = stats.skew(daily_base)
```

```

excess_kurt = stats.kurtosis(daily_base) # Fisher definition (excess)
sr_se = np.sqrt(
    (1 + 0.5 * sr**2 - skewness * sr + ((excess_kurt + 2) / 4) * sr**2)
    / (n - 1)
)
sr_ci_lo = sr - 1.96 * sr_se
sr_ci_hi = sr + 1.96 * sr_se

print("=== 16b: Sharpe ratio confidence interval (Lo 2002) ===")
print(f" Annualised Sharpe : {sr:.3f}")
print(f" Skewness          : {skewness:.3f}")
print(f" Excess kurtosis      : {excess_kurt:.3f}")
print(f" SE of Sharpe         : {sr_se:.3f}")
print(f" 95% CI              : [{sr_ci_lo:.3f}, {sr_ci_hi:.3f}]")
print()
if sr_ci_lo > 0:
    print(" Interpretation: the entire 95% CI lies above zero.")
    print(" The annualised Sharpe is significantly positive at t+0.")
else:
    print(" Interpretation: CI crosses zero -- Sharpe is not reliably positive.")

```

```

=== 16a: One-sample t-test -- daily return vs 0 ===
n          : 159
mean return : 0.0099 (0.99% per day)
t-statistic : 2.641
p-value     : 0.0091
95% CI      : [0.0026, 0.0172]

```

Interpretation: reject H0 at 5% -- mean return is significantly positive.

```

=== 16b: Sharpe ratio confidence interval (Lo 2002) ===
Annualised Sharpe : 3.325
Skewness          : 0.188
Excess kurtosis    : 1.221
SE of Sharpe       : 0.306
95% CI            : [2.725, 3.925]

```

Interpretation: the entire 95% CI lies above zero.
The annualised Sharpe is significantly positive at t+0.

17. Walk-Forward Robustness Test

A single in-sample Sharpe ratio is insufficient evidence of a persistent edge. We split the full date range into three equal, non-overlapping folds and measure performance independently in each.

If the signal is real, we expect positive Sharpe in all three folds. Deterioration over time is consistent with a decaying microstructure edge -- still meaningful, but worth monitoring.

Scenario evaluated: t+0 execution, 20 bps transaction cost.

```

In [18]: # --- Walk-forward: 3 equal folds ---
daily_base = scenario_daily(event_filtered_df, delay=0, cost_bps=20)
dates_idx = daily_base.index
n_dates = len(dates_idx)
fold_size = n_dates // 3

wf_results = []

print("=== Walk-forward: 3-fold split (t+0, 20 bps) ===")
print(f" Total trading days : {n_dates}")
print(f" Days per fold       : ~{fold_size}")
print()

for i in range(3):
    start = dates_idx[i * fold_size]
    end = dates_idx[min((i + 1) * fold_size - 1, n_dates - 1)]
    fold = daily_base.loc[start:end].dropna()

    if len(fold) < 2:
        continue

    fold_sr = (fold.mean() / fold.std()) * np.sqrt(trading_days)
    fold_ret = (1 + fold).prod() - 1
    fold_eq = (1 + fold).cumprod()
    fold_dd = (fold_eq / fold_eq.cummax() - 1).min()

    wf_results.append({
        "fold": i + 1,
        "start": start.date(),
        "end": end.date(),
        "n": len(fold),
        "sharpe": round(fold_sr, 2),
        "total_return": round(fold_ret, 3),
        "max_drawdown": round(fold_dd, 3),
    })

    print(f" Fold {i+1}: {start.date()} -> {end.date()}")
    print(f" n          : {len(fold)}")
    print(f" Sharpe     : {fold_sr:.2f}")
    print(f" Total return : {fold_ret:.1%}")
    print(f" Max drawdown : {fold_dd:.1%}")
    print()

wf_df = pd.DataFrame(wf_results)

print("Summary:")
print(wf_df.to_string(index=False))
print()

sharpes = wf_df["sharpe"]
print(f" Mean fold Sharpe : {sharpes.mean():.2f}")
print(f" Min fold Sharpe : {sharpes.min():.2f}")
print()

```

```
if (sharpes > 0).all():
    print(" Interpretation: Sharpe is positive in all three folds.")
    if sharpes.iloc[0] > sharpes.iloc[-1]:
        print(" Note: Sharpe declines across folds, consistent with a decaying microstructure edge.")
    else:
        print(" Sharpe is stable or improving across folds.")
else:
    print(" Warning: at least one fold has negative Sharpe -- signal is not consistent.")

# persist walk-forward results
fp = BACKTEST_DIR / "walkforward_results.pkl"
wf_df.to_pickle(fp)
manifest[stage_label]["walkforward_results"] = str(fp.relative_to(OUTPUT_ROOT))
print(f"\n saved: {fp}")
```

=== Walk-forward: 3-fold split (t+0, 20 bps) ===

Total trading days : 159
Days per fold : ~53

Fold 1: 2023-08-04 -> 2024-05-22

n : 53
Sharpe : 4.26
Total return : 86.0%
Max drawdown : -14.2%

Fold 2: 2024-06-05 -> 2025-02-05

n : 53
Sharpe : 3.85
Total return : 75.1%
Max drawdown : -13.2%

Fold 3: 2025-02-06 -> 2026-05-13

n : 53
Sharpe : 1.75
Total return : 23.7%
Max drawdown : -22.1%

Summary:

fold	start	end	n	sharpe	total_return	max_drawdown
1	2023-08-04	2024-05-22	53	4.26	0.860	-0.142
2	2024-06-05	2025-02-05	53	3.85	0.751	-0.132
3	2025-02-06	2026-05-13	53	1.75	0.237	-0.221

Mean fold Sharpe : 3.29
Min fold Sharpe : 1.75

Interpretation: Sharpe is positive in all three folds.
Note: Sharpe declines across folds, consistent with a decaying microstructure edge.

saved: ..\output\005_backtest\walkforward_results.pkl

999. Persist the Manifest

Records all artifact paths and a completion timestamp for this stage into the shared manifest.

```
In [19]: manifest[stage_label]["timestamp"] = datetime.now(UTC).isoformat()
pd.to_pickle(manifest, MANIFEST_FILE)
print("manifest saved:", MANIFEST_FILE)
print("pipeline stages:", sorted(manifest.keys()))
print(f"artifacts in {stage_label}:", sorted(manifest[stage_label].keys()))
```

manifest saved: ..\output\manifest.pkl
pipeline stages: ['001_download', '002_enrich', '003_analysis', '003a_regimes', '004_strategy', '005_backtest']
artifacts in 005_backtest: ['daily_returns', 'equity_curve', 'grid_results', 'performance_summary', 'timestamp', 'trade_df', 'walkforward_results']

In []: